

Selectively instrumenting the code with DR_HOOK

Philippe Marguinaud, 22 June 2010

The whole ARPEGE/AROME code (including the XRD, and SURFEX) has been instrumented with `dr_hook` so that it be possible to assess the model performances on several platforms without relying on proprietary instrumentation software.

We found that running AROME on the NEC with `dr_hook` enabled yielded unreasonable overhead (more than 80%; this is the assessment of `dr_hook` itself).

Therefore, re-compiled the code with `gmckpack` using the `GMK_DR_HOOK_ALL` option. This option allows to select which files will be instrumented at run-time, using a namelist.

This namelist can be created automatically from existing `dr_hook` profiles using the following scripts:

- `mkdh2f.pl`; scans the code to establish the mapping between `dr_hook` tags and source code files; it should be run from the pack directory being considered and generates a file named `'dh2f.pl'`
- `mknamhook.pl`; takes a `dr_hook` profile, a threshold in milliseconds, and the name of file in which the namelist containing the source files not to be instrumented will be written. Routines whose elapsed an CPU time are below this threshold will not be evaluated by `dr_hook`.

Note that it is not always possible to remove a file from `dr_hook` instrumentation because a single source file may contain several routines with different profiles; a conservative approach is used : a source file is kept if at least one of its routine has either an average elapsed or CPU time above the threshold passed as argument. The list of source file names which should have been removed but could not is printed by `mknamhook.pl`.

After filtering files with `mknamhook.pl`, the overhead was reduced to about 8%; it is worth mentioning that `mknamhook.pl` showed that some conflicts remained (some source files contain several subroutines with different profiles, that is with times above and below the threshold). An overhead of 8% can however be considered as acceptable.

Below is the approach we recommend for using `dr_hook`:

- compile your pack with the “`export GMK_DR_HOOK_ALL=1`” option
- run the software once with `dr_hook` enabled on all routines
- filter the routines whose CPU or elapsed time is too small
- run the software again; the `dr_hook` overhead should be dramatically reduced

One issue is to find the good threshold to pass as an argument to `mknamhook.pl`. Here is how we decided to use 3 ms on the NEC; for this purpose, we used the following program:

```
program main
use yomhook, only : lhook, dr_hook
```

```

implicit none
integer :: i
real*8 :: zhook_handle
if (lhook) call dr_hook('main', 0, zhook_handle)
do i = 1, 100000
    call sub
enddo
if (lhook) call dr_hook('main', 1, zhook_handle)
contains
subroutine sub
real*8 :: zhook_handle
if (lhook) call dr_hook('sub',0,zhook_handle)
if (lhook) call dr_hook('sub',1,zhook_handle)
end subroutine
end program

```

On the NEC, running this program with dr_hook took about 32s, and without dr_hook, it took 3s; the overhead of dr_hook is therefore 0.3 ms per routine. In order to have an overhead smaller than 10%, it is therefore necessary to remove subroutines whose CPU or elapsed time is below 3 ms.